Module Review

1. What are the common big data challenges that you will be building solutions for in this course? (check all that apply)

- **Migrating existing on-premise workloads to the cloud**
- **Analyzing large datasets at scale**
- Building containerized applications for web development
- **Building streaming data pipelines**
- **Applying machine learning to your datasets**

2. You have a large enterprise that will likely have many teams using their own Google Cloud Platform projects and resources. What should you be sure to have to help manage and administer these resources? (check all that apply)

- **A defined Organization**
- **Folders for teams and/or products**
- **A defined access control policy with Cloud IAM**
- A Kubernetes or Hadoop cluster for each project

3. Which of the following is **NOT** one of the advantages of Google Cloud security

- **Google Cloud will automatically manage and curate your content and access policies to be safe for the public**
- Google Cloud will secure the physical hardware that is running your applications and infrastructure
- Google Cloud has tools like Cloud IAM that help you administer and set company-wide security policies
- Google Cloud will manage audit logging of access and use of resources in your account

4. If you don't have a large dataset of your own but still want to practice writing queries and building pipelines on Google Cloud Platform, what should you do?

- Practice with the datasets in the Google Cloud Public Datasets program
- **Find other public datasets online and upload them into BigQuery**
- Work to create your own dataset and then upload it into BigQuery for analysis

5. As you saw in the demo, Compute Engine nodes on GCP are:

- Expensive to create and teardown
- One of ~50 choices in terms of CPU and memory
- Pre-installed with all the software packages you might ever need.
- **Allocated on demand, and you pay for the time that they are up.**